

activities_analysis

August 21, 2024

1 Strava Running Data Analysis: Performance Trends and Influencing Factors

1.1 Foreword

“Pain is inevitable. Suffering is optional.”

A quote from one of the only books I had read about fitness. "What I Talk About When I Talk About Running" written by Haruki Murakami. I was curious what was a man like this doing in the fitness section. At that point I hadn't run a single yard more than what was required of me – but by the time I finished reading, I was a changed man.

I slowly built a habit of running from 2016, enjoyed the meditative aspect of it, ran often enough but never really got out of the proverbial shallow end of the pool. Its now 2024 and after many fun years of exploring other sports, I realized my base fitness was lacking. I had decided to pick up running sometime in May, for real this time. Now I've signed up for a half marathon in November, and I want to use what I know about data science to assist me.

1.2 Abstract

In this notebook, we will dive in and analyze running data collected from my *Strava* to identify performance trends and influencing factors. The analysis will not only provide insights into overall fitness progression but also inform on the factors that influence running performance as well as the distinct types of runs that are most common.

1.3 Project steps

1. Source data from Strava Export. Download files from [Bulk Export](#) on my own profile menu.
2. Load, assess and clean the data for analysis.
3. Create any features that would provide additional insight for analysis.
4. Generate graphs and charts to illustrate insights.
5. Analyze with machine learning.

1.4 Table of Contents

- Accessing the data;
- Data Cleaning;
- Feature Engineering;
- **Insights**;
- Feature Selection;

- Machine Learning;
- Conclusion;

```
[1]: #pandas and plotting
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import matplotlib.dates as mdates
import calendar
from datetime import timedelta
from matplotlib.patches import FancyArrowPatch

# machine learning
from sklearn import preprocessing
from sklearn import metrics
from sklearn.cluster import KMeans
from sklearn.feature_selection import chi2
from sklearn.feature_selection import f_regression
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.model_selection import train_test_split
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler

%config IPCompleter.greedy=True
import warnings
warnings.filterwarnings('ignore')
```

2 Accessing the data

```
[2]: # Load the data from an export on Strava
original = pd.read_csv('../Data/activities.csv')
print('Dataframe shape:', original.shape)
original.head()
```

Dataframe shape: (254, 94)

```
[2]:
```

	Activity ID	Activity Date	Activity Name	Activity Type	\
0	676696158	Aug 15, 2016, 9:53:13 AM	Afternoon Run	Run	
1	681425706	Aug 19, 2016, 10:11:34 AM	Evening Run	Run	
2	683559635	Aug 21, 2016, 10:14:46 AM	Evening Run	Run	

3	689475888	Aug 26, 2016, 10:22:15 AM	Evening Run	Run
4	708051143	Sep 11, 2016, 9:19:10 AM	Afternoon Run	Run

	Activity Description	Elapsed Time	Distance	Max Heart Rate	\
0	NaN	1414	3.31	NaN	
1	NaN	2062	4.45	NaN	
2	NaN	1903	4.41	NaN	
3	NaN	1920	4.51	NaN	
4	NaN	1922	4.43	NaN	

	Relative Effort	Commute	...	Activity Count	Total Steps	Carbon Saved	\
0	NaN	False	...	NaN	NaN	NaN	
1	NaN	False	...	NaN	NaN	NaN	
2	NaN	False	...	NaN	NaN	NaN	
3	NaN	False	...	NaN	NaN	NaN	
4	NaN	False	...	NaN	NaN	NaN	

	Pool Length	Training Load	Intensity	Average Grade	Adjusted Pace	\
0	NaN	NaN	NaN		NaN	
1	NaN	NaN	NaN		NaN	
2	NaN	NaN	NaN		NaN	
3	NaN	NaN	NaN		NaN	
4	NaN	NaN	NaN		NaN	

	Timer Time	Total Cycles	Media
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN

[5 rows x 94 columns]

2.0.1 First observations:

There are a whopping 94 columns, most of which are either null or dont contribute to the analysis, so lets keep only the relevant columns for research

3 Data cleaning

```
[3]: original.columns = original.columns.str.lower().str.replace(' ','_') #
      ↪Standardize the column titles
original = original[original['activity_type'] == 'Run'] # Filter to only
      ↪Running values
df = original.copy(deep=True) # Make a copy so that there is always the
      ↪original to reference to
```

```
[4]: print(df.shape)
df.head()
```

(217, 94)

```
[4]: activity_id          activity_date activity_name activity_type \
0    676696158    Aug 15, 2016, 9:53:13 AM Afternoon Run      Run
1    681425706    Aug 19, 2016, 10:11:34 AM  Evening Run      Run
2    683559635    Aug 21, 2016, 10:14:46 AM  Evening Run      Run
3    689475888    Aug 26, 2016, 10:22:15 AM  Evening Run      Run
4    708051143    Sep 11, 2016, 9:19:10 AM  Afternoon Run      Run

activity_description elapsed_time distance max_heart_rate \
0                NaN         1414      3.31          NaN
1                NaN         2062      4.45          NaN
2                NaN         1903      4.41          NaN
3                NaN         1920      4.51          NaN
4                NaN         1922      4.43          NaN

relative_effort  commute  ... activity_count total_steps carbon_saved \
0                NaN    False  ...          NaN          NaN          NaN
1                NaN    False  ...          NaN          NaN          NaN
2                NaN    False  ...          NaN          NaN          NaN
3                NaN    False  ...          NaN          NaN          NaN
4                NaN    False  ...          NaN          NaN          NaN

pool_length  training_load  intensity  average_grade_adjusted_pace \
0            NaN          NaN        NaN          NaN
1            NaN          NaN        NaN          NaN
2            NaN          NaN        NaN          NaN
3            NaN          NaN        NaN          NaN
4            NaN          NaN        NaN          NaN

timer_time  total_cycles  media
0          NaN          NaN    NaN
1          NaN          NaN    NaN
2          NaN          NaN    NaN
3          NaN          NaN    NaN
4          NaN          NaN    NaN
```

[5 rows x 94 columns]

```
[5]: def list_nulls(df):
    null_df = [[col, df[col].isnull().sum()] for col in df.columns]
    print('Null Data:', df.isnull().sum().sum())
    print(list(filter(lambda x: x[1]>0, null_df)))
```

```
list_nulls(df)
```

Null Data: 14355

```
[['activity_description', 187], ['max_heart_rate', 160], ['relative_effort', 160], ['activity_private_note', 217], ['activity_gear', 168], ['athlete_weight', 92], ['bike_weight', 217], ['average_speed', 90], ['elevation_loss', 100], ['average_positive_grade', 217], ['average_negative_grade', 217], ['max_cadence', 217], ['average_cadence', 217], ['max_heart_rate.1', 170], ['average_heart_rate', 160], ['max_watts', 217], ['average_watts', 217], ['max_temperature', 217], ['average_temperature', 217], ['relative_effort.1', 160], ['total_work', 217], ['number_of_runs', 217], ['uphill_time', 217], ['downhill_time', 217], ['other_time', 217], ['perceived_exertion', 215], ['type', 217], ['start_time', 217], ['weighted_average_power', 217], ['power_count', 217], ['prefer_perceived_exertion', 185], ['perceived_relative_effort', 215], ['total_weight_lifted', 217], ['from_upload', 90], ['grade_adjusted_distance', 62], ['weather_observation_time', 190], ['weather_condition', 190], ['weather_temperature', 190], ['apparent_temperature', 190], ['dewpoint', 190], ['humidity', 190], ['weather_pressure', 190], ['wind_speed', 190], ['wind_gust', 190], ['wind_bearing', 190], ['precipitation_intensity', 190], ['sunrise_time', 190], ['sunset_time', 190], ['moon_phase', 190], ['bike', 217], ['gear', 166], ['precipitation_probability', 190], ['precipitation_type', 191], ['cloud_cover', 190], ['weather_visibility', 190], ['uv_index', 190], ['weather_ozone', 209], ['jump_count', 217], ['total_grit', 217], ['average_flow', 217], ['flagged', 155], ['average_elapsed_speed', 155], ['dirt_distance', 154], ['newly_explored_distance', 217], ['newly_explored_dirt_distance', 217], ['activity_count', 217], ['total_steps', 155], ['carbon_saved', 217], ['pool_length', 217], ['training_load', 217], ['intensity', 217], ['average_grade_adjusted_pace', 176], ['timer_time', 217], ['total_cycles', 217], ['media', 199]]
```

```
[6]: columns_keep = [
    ['activity_date', 'moving_time', 'distance', 'max_speed', 'average_speed', 'elevation_gain', 'elevation_loss']
df = df.loc[:, columns_keep]
print(df.shape)
df.head()
```

(217, 14)

```
[6]:
```

	activity_date	moving_time	distance	max_speed	average_speed	\
0	Aug 15, 2016, 9:53:13 AM	1314.0	3.31	7.2	NaN	
1	Aug 19, 2016, 10:11:34 AM	1802.0	4.45	8.3	NaN	
2	Aug 21, 2016, 10:14:46 AM	1898.0	4.41	5.5	NaN	
3	Aug 26, 2016, 10:22:15 AM	1900.0	4.51	5.7	NaN	
4	Sep 11, 2016, 9:19:10 AM	1913.0	4.43	4.9	NaN	

	elevation_gain	elevation_loss	elevation_low	elevation_high	max_grade	\
0	39.465000	NaN	53.200001	87.599998	16.200001	
1	63.403999	NaN	50.299999	87.300003	40.200001	
2	51.626202	NaN	46.400002	87.699997	16.299999	

3	119.328003	NaN	64.000000	137.699997	22.400000
4	130.173996	NaN	64.000000	138.300003	23.500000

	average_grade	max_heart_rate	average_heart_rate	total_steps
0	-0.241473	NaN	NaN	NaN
1	0.125682	NaN	NaN	NaN
2	-0.002268	NaN	NaN	NaN
3	0.026587	NaN	NaN	NaN
4	0.069921	NaN	NaN	NaN

```
[7]: df = df.dropna(axis=1,how='all') # Dropping all the rows with completely null
      ↪rows
      df.shape
```

[7]: (217, 14)

```
[8]: # Convert 'activity_date' to datetime
      df['activity_date'] = pd.to_datetime(df['activity_date'])
```

```
[9]: df = df.drop(df[(df.distance < 1)].index) # Remove any runs under 1km, most
      ↪likely misinputs or warmups
      df = df.reset_index(drop=True)
      print(df.shape)
      df.head()
```

(212, 14)

```
[9]:
```

	activity_date	moving_time	distance	max_speed	average_speed	\
0	2016-08-15 09:53:13	1314.0	3.31	7.2	NaN	
1	2016-08-19 10:11:34	1802.0	4.45	8.3	NaN	
2	2016-08-21 10:14:46	1898.0	4.41	5.5	NaN	
3	2016-08-26 10:22:15	1900.0	4.51	5.7	NaN	
4	2016-09-11 09:19:10	1913.0	4.43	4.9	NaN	

	elevation_gain	elevation_loss	elevation_low	elevation_high	max_grade	\
0	39.465000	NaN	53.200001	87.599998	16.200001	
1	63.403999	NaN	50.299999	87.300003	40.200001	
2	51.626202	NaN	46.400002	87.699997	16.299999	
3	119.328003	NaN	64.000000	137.699997	22.400000	
4	130.173996	NaN	64.000000	138.300003	23.500000	

	average_grade	max_heart_rate	average_heart_rate	total_steps
0	-0.241473	NaN	NaN	NaN
1	0.125682	NaN	NaN	NaN
2	-0.002268	NaN	NaN	NaN
3	0.026587	NaN	NaN	NaN
4	0.069921	NaN	NaN	NaN

```
[10]: # Check for nulls again, consider any imputes or further drops
list_nulls(df)
```

Null Data: 665

```
[['average_speed', 90], ['elevation_loss', 100], ['max_heart_rate', 160],
['average_heart_rate', 160], ['total_steps', 155]]
```

4 Feature engineering

```
[11]: # Create year and month columns
df['year'] = df['activity_date'].dt.year
df['month'] = df['activity_date'].dt.month

#Convert the speeds into a more familiar metric which is minutes/km
df['moving_time_minutes'] = round(df['moving_time']/60, 2)
df['distance'] = round(df['distance'], 2) # Since its already in km, no need
↳for any conversion
df['pace'] = round(df['moving_time_minutes'] / df['distance'],2)
df['max_pace'] = round(1000/df['max_speed'] / 60,2)
df = df.drop(columns=['average_speed', 'max_speed'])

# Add some useful date & time features
df['week'] = df['activity_date'].dt.to_period('W')
df['start_date'] = df['activity_date'].dt.date
df['start_time'] = df['activity_date'].dt.time
df['week_start_date'] = df.week.apply(lambda r: r.start_time.date())
df['weekday'] = df['activity_date'].apply(lambda x: x.weekday())
df.head()
```

```
[11]:      activity_date  moving_time  distance  elevation_gain  elevation_loss \
0 2016-08-15 09:53:13      1314.0      3.31      39.465000      NaN
1 2016-08-19 10:11:34      1802.0      4.45      63.403999      NaN
2 2016-08-21 10:14:46      1898.0      4.41      51.626202      NaN
3 2016-08-26 10:22:15      1900.0      4.51      119.328003      NaN
4 2016-09-11 09:19:10      1913.0      4.43      130.173996      NaN
```

```
      elevation_low  elevation_high  max_grade  average_grade  max_heart_rate \
0      53.200001      87.599998  16.200001      -0.241473      NaN
1      50.299999      87.300003  40.200001      0.125682      NaN
2      46.400002      87.699997  16.299999      -0.002268      NaN
3      64.000000      137.699997  22.400000      0.026587      NaN
4      64.000000      138.300003  23.500000      0.069921      NaN
```

```
      ...  year  month  moving_time_minutes  pace  max_pace \
0  ...  2016      8      21.90  6.62      2.31
1  ...  2016      8      30.03  6.75      2.01
2  ...  2016      8      31.63  7.17      3.03
```

3	...	2016	8	31.67	7.02	2.92
4	...	2016	9	31.88	7.20	3.40

		week	start_date	start_time	week_start_date	weekday
0	2016-08-15/2016-08-21		2016-08-15	09:53:13	2016-08-15	0
1	2016-08-15/2016-08-21		2016-08-19	10:11:34	2016-08-15	4
2	2016-08-15/2016-08-21		2016-08-21	10:14:46	2016-08-15	6
3	2016-08-22/2016-08-28		2016-08-26	10:22:15	2016-08-22	4
4	2016-09-05/2016-09-11		2016-09-11	09:19:10	2016-09-05	6

[5 rows x 22 columns]

4.0.1 –Excess of nulls in the columns Elevation loss, Max Heart Rate, Average Heart Rate and Total Steps–

Besides pace, I suspect the other nulls are due to the fact that majority of my runs were run **before** I bought a fitness tracker watch. Unfortunately, will have to drop them.

But instead of dropping the columns, I will instead *filter the dataframe* to the time I started training in 2024, which should have a complete set of data.

```
[12]: latest_df = df[df.activity_date > '2024-04-01']
list_nulls(latest_df)
print(latest_df.shape)
latest_df.head()
```

Null Data: 2

```
[['max_heart_rate', 1], ['average_heart_rate', 1]]
(43, 22)
```

```
[12]:
```

	activity_date	moving_time	distance	elevation_gain	\
169	2024-04-23 03:41:19	1005.0	1.36	0.000000	
170	2024-04-24 00:04:13	2312.0	4.05	168.199997	
171	2024-04-25 01:15:43	1117.0	1.35	0.000000	
172	2024-05-01 08:12:56	2280.0	4.07	278.000000	
173	2024-04-29 03:30:40	1963.0	2.01	0.000000	

	elevation_loss	elevation_low	elevation_high	max_grade	average_grade	\
169	0.000000	0.000000	0.000000	0.000000	0.000000	
170	166.199997	36.400002	104.599998	48.522873	0.049375	
171	0.000000	0.000000	0.000000	0.000000	0.000000	
172	282.799988	72.199997	126.599998	49.813019	-0.117833	
173	0.000000	0.000000	0.000000	0.000000	0.000000	

	max_heart_rate	...	year	month	moving_time_minutes	pace	max_pace	\
169	186.0	...	2024	4	16.75	12.32	6.01	
170	174.0	...	2024	4	38.53	9.51	4.48	
171	175.0	...	2024	4	18.62	13.79	5.36	
172	174.0	...	2024	5	38.00	9.34	3.95	

173	183.0	...	2024	4	32.72	16.28	6.56
-----	-------	-----	------	---	-------	-------	------

	week	start_date	start_time	week_start_date	weekday
169	2024-04-22/2024-04-28	2024-04-23	03:41:19	2024-04-22	1
170	2024-04-22/2024-04-28	2024-04-24	00:04:13	2024-04-22	2
171	2024-04-22/2024-04-28	2024-04-25	01:15:43	2024-04-22	3
172	2024-04-29/2024-05-05	2024-05-01	08:12:56	2024-04-29	2
173	2024-04-29/2024-05-05	2024-04-29	03:30:40	2024-04-29	0

[5 rows x 22 columns]

```
[13]: # Impute the remaining missing data with the averages
latest_df['average_heart_rate'] = latest_df['average_heart_rate'].
    ↪fillna(value=latest_df['average_heart_rate'].mean())
latest_df['max_heart_rate'] = latest_df['max_heart_rate'].
    ↪fillna(value=latest_df['max_heart_rate'].mean())
list_nulls(latest_df)
```

Null Data: 0

[]

```
[14]: zero_rows = latest_df[(latest_df == 0).any(axis=1)]
print(zero_rows.shape)
zero_rows
```

(8, 22)

```
[14]:      activity_date  moving_time  distance  elevation_gain  \
169 2024-04-23 03:41:19      1005.0      1.36          0.0
171 2024-04-25 01:15:43      1117.0      1.35          0.0
173 2024-04-29 03:30:40      1963.0      2.01          0.0
175 2024-05-07 02:21:57      1714.0      3.28          0.0
176 2024-05-08 03:55:46      2147.0      2.09          0.0
184 2024-06-24 22:34:25       2023.0      3.87        200.0
199 2024-07-24 14:49:33      1086.0      2.24          0.0
206 2024-08-05 23:19:57      1254.0      2.18        143.0

      elevation_loss  elevation_low  elevation_high  max_grade  average_grade  \
169      0.000000      0.0      0.000000      0.000000      0.000000
171      0.000000      0.0      0.000000      0.000000      0.000000
173      0.000000      0.0      0.000000      0.000000      0.000000
175      0.000000      0.0      0.000000      0.000000      0.000000
176      0.000000      0.0      0.000000      0.000000      0.000000
184      219.800003      15.6      85.000000      48.284103      -0.517903
199      0.000000      0.0      0.000000      0.000000      0.000000
206      140.000000      16.4      63.200001      48.390980      0.099630

      max_heart_rate  ...  year  month  moving_time_minutes  pace  max_pace  \
```

169	186.0	...	2024	4	16.75	12.32	6.01
171	175.0	...	2024	4	18.62	13.79	5.36
173	183.0	...	2024	4	32.72	16.28	6.56
175	189.0	...	2024	5	28.57	8.71	6.46
176	190.0	...	2024	5	35.78	17.12	6.66
184	168.0	...	2024	6	33.72	8.71	3.79
199	180.0	...	2024	7	18.10	8.08	6.35
206	153.0	...	2024	8	20.90	9.59	4.88

	week	start_date	start_time	week_start_date	weekday
169	2024-04-22/2024-04-28	2024-04-23	03:41:19	2024-04-22	1
171	2024-04-22/2024-04-28	2024-04-25	01:15:43	2024-04-22	3
173	2024-04-29/2024-05-05	2024-04-29	03:30:40	2024-04-29	0
175	2024-05-06/2024-05-12	2024-05-07	02:21:57	2024-05-06	1
176	2024-05-06/2024-05-12	2024-05-08	03:55:46	2024-05-06	2
184	2024-06-24/2024-06-30	2024-06-24	22:34:25	2024-06-24	0
199	2024-07-22/2024-07-28	2024-07-24	14:49:33	2024-07-22	2
206	2024-08-05/2024-08-11	2024-08-05	23:19:57	2024-08-05	0

[8 rows x 22 columns]

4.0.2 –Zero values in many elevation metrics–

With a total of 6 rows containing 0 elevation data, it is worth considering dropping these rows.

Checking on actual strava, I found that these were manual uploads that did not upload properly. Will drop these.

```
[15]: latest_df = latest_df.drop(zero_rows.index)
latest_df.shape
```

```
[15]: (35, 22)
```

```
[16]: latest_df.head()
```

```
[16]:
```

	activity_date	moving_time	distance	elevation_gain	\
170	2024-04-24 00:04:13	2312.0	4.05	168.199997	
172	2024-05-01 08:12:56	2280.0	4.07	278.000000	
174	2024-05-04 09:41:57	2359.0	5.01	193.399994	
177	2024-05-17 22:25:45	1937.0	3.40	166.400024	
178	2024-05-18 23:22:30	8935.0	14.81	992.799927	

	elevation_loss	elevation_low	elevation_high	max_grade	average_grade	\
170	166.199997	36.400002	104.599998	48.522873	0.049375	
172	282.799988	72.199997	126.599998	49.813019	-0.117833	
174	164.600006	-1.600000	37.799999	49.948692	0.570642	
177	166.399994	442.399994	481.600006	48.871666	0.047297	
178	1008.799988	395.200012	523.000000	49.956367	-0.156001	

	max_heart_rate	...	year	month	moving_time_minutes	pace	max_pace	\
170	174.0	...	2024	4	38.53	9.51	4.48	
172	174.0	...	2024	5	38.00	9.34	3.95	
174	182.0	...	2024	5	39.32	7.85	1.23	
177	166.0	...	2024	5	32.28	9.49	4.27	
178	202.0	...	2024	5	148.92	10.06	4.98	

	week	start_date	start_time	week_start_date	weekday
170	2024-04-22/2024-04-28	2024-04-24	00:04:13	2024-04-22	2
172	2024-04-29/2024-05-05	2024-05-01	08:12:56	2024-04-29	2
174	2024-04-29/2024-05-05	2024-05-04	09:41:57	2024-04-29	5
177	2024-05-13/2024-05-19	2024-05-17	22:25:45	2024-05-13	4
178	2024-05-13/2024-05-19	2024-05-18	23:22:30	2024-05-13	5

[5 rows x 22 columns]


5 Insights

Now we will be exploring some insights using data visualization.

Note: pace is in minutes/km, therefore the *lower* the y-value, the *faster* it is. So lower=better.


```
[17]: sns.set(style='ticks',context='talk')
sns.regplot(x='distance',y='max_pace',data=latest_df).set_title('Max pace vs_
↳distance')
```

```
[17]: Text(0.5, 1.0, 'Max pace vs distance')
```




```
[18]: sns.set(style='ticks',context='talk')
sns.regplot(x='distance',y='pace',data=latest_df).set_title('Pace vs distance')
```

```
[18]: Text(0.5, 1.0, 'Pace vs distance')
```



```
[19]: # Calculate rolling average pace to smooth out variations
rolling = latest_df.pace.rolling(window=5).mean()

plt.figure(figsize=(10, 6))
plt.plot(latest_df.start_date, rolling)
plt.title('Rolling Average Pace over Time')
plt.xlabel('Date')
plt.ylabel('Pace (minutes/km)')
plt.show()
```



5.0.1 Rolling Average Insights

The data exhibits a recurring pattern of sharp performance improvements followed by notable pace increases, likely indicative of intense training phases succeeded by recovery periods. Significantly, this cyclical pattern demonstrates a gradual attenuation over time, suggesting:

- Enhanced recovery efficiency
- Improved baseline fitness
- Potential optimization of the training-recovery balance

This trend points towards a maturing and increasingly effective training regimen, where the performance gains more consistently while potentially reducing the physiological stress of intense training phases.

```
[20]: class BubbleChart:
    def __init__(self, area, bubble_spacing=0):
        area = np.asarray(area)
        r = np.sqrt(area / np.pi)
        self.bubble_spacing = bubble_spacing
        self.bubbles = np.ones((len(area), 4))
        self.bubbles[:, 2] = r
        self.bubbles[:, 3] = area
        self.maxstep = 2 * self.bubbles[:, 2].max() + self.bubble_spacing
        self.step_dist = self.maxstep / 2
        length = np.ceil(np.sqrt(len(self.bubbles)))
```

```

        grid = np.arange(length) * self.maxstep
        gx, gy = np.meshgrid(grid, grid)
        self.bubbles[:, 0] = gx.flatten()[:len(self.bubbles)]
        self.bubbles[:, 1] = gy.flatten()[:len(self.bubbles)]
        self.com = self.center_of_mass()

    def center_of_mass(self):
        return np.average(self.bubbles[:, :2], axis=0, weights=self.bubbles[:, 2])

    def center_distance(self, bubble, bubbles):
        return np.hypot(bubble[0] - bubbles[:, 0], bubble[1] - bubbles[:, 1])

    def outline_distance(self, bubble, bubbles):
        center_distance = self.center_distance(bubble, bubbles)
        return center_distance - bubble[2] - bubbles[:, 2] - self.bubble_spacing

    def check_collisions(self, bubble, bubbles):
        distance = self.outline_distance(bubble, bubbles)
        return len(distance[distance < 0])

    def collides_with(self, bubble, bubbles):
        distance = self.outline_distance(bubble, bubbles)
        idx_min = np.argmin(distance)
        return idx_min if type(idx_min) == np.ndarray else [idx_min]

    def collapse(self, n_iterations=50):
        for _i in range(n_iterations):
            moves = 0
            for i in range(len(self.bubbles)):
                rest_bub = np.delete(self.bubbles, i, 0)
                dir_vec = self.com - self.bubbles[i, :2]
                dir_vec = dir_vec / np.sqrt(dir_vec.dot(dir_vec))
                new_point = self.bubbles[i, :2] + dir_vec * self.step_dist
                new_bubble = np.append(new_point, self.bubbles[i, 2:4])
                if not self.check_collisions(new_bubble, rest_bub):
                    self.bubbles[i, :] = new_bubble
                    self.com = self.center_of_mass()
                    moves += 1
            else:
                for colliding in self.collides_with(new_bubble, rest_bub):
                    dir_vec = rest_bub[colliding, :2] - self.bubbles[i, :2]
                    dir_vec = dir_vec / np.sqrt(dir_vec.dot(dir_vec))
                    orth = np.array([dir_vec[1], -dir_vec[0]])
                    new_point1 = self.bubbles[i, :2] + orth * self.step_dist
                    new_point2 = self.bubbles[i, :2] - orth * self.step_dist

```

```

        dist1 = self.center_distance(self.com, np.
↪array([new_point1]))
        dist2 = self.center_distance(self.com, np.
↪array([new_point2]))
        new_point = new_point1 if dist1 < dist2 else new_point2
        new_bubble = np.append(new_point, self.bubbles[i, 2:4])
        if not self.check_collisions(new_bubble, rest_bub):
            self.bubbles[i, :] = new_bubble
            self.com = self.center_of_mass()
        if moves / len(self.bubbles) < 0.1:
            self.step_dist = self.step_dist / 2

def plot(self, ax, data, colors):
    for i, (year, total, avg) in enumerate(data):
        bubble = self.bubbles[i]

        # Create the bubble
        circ = plt.Circle(bubble[:2], bubble[2], color=colors[i], alpha=0.7)
        ax.add_patch(circ)

        # Add year in the center
        ax.text(*bubble[:2], str(year),
                horizontalalignment='center', verticalalignment='center',
                color='black', fontname='Arial', fontsize=12,
↪fontweight='bold')

        # Add single arrow with both annotations
        self._add_annotations(ax, bubble, total, avg)

def _add_annotations(self, ax, bubble, total, avg):
    radius = bubble[2]
    angle = 45 # You can adjust this angle as needed

    # Calculate the position for the arrow and text
    arrow_x = bubble[0] + np.cos(np.radians(angle)) * radius * 0.7
    arrow_y = bubble[1] + np.sin(np.radians(angle)) * radius * 0.7

    text_x = arrow_x + np.cos(np.radians(angle)) # * radius * 0.5
    text_y = arrow_y + np.sin(np.radians(angle)) # * radius * 0.5

    # Create the arrow
    ax.annotate('', xy=(bubble[0] + 0.5, bubble[1] + 0.5), xytext=(arrow_x,
↪arrow_y),
    ↪arrowprops=dict(arrowstyle="->",
↪connectionstyle="arc3,rad=0.2", color='black'))

    # Add the text annotations

```



```

        ax.text(text_x, text_y, f"Total: {total:.0f} km\nAvg: {avg:.2f} km",
                horizontalalignment='left', verticalalignment='center',
                color='black', fontname='Arial', fontsize=8,
                bbox=dict(facecolor='white', edgecolor='gray', alpha=0.7,
        pad=4))

```

```

[21]: # Prepare data for bubble chart
yearly_stats = df.groupby('year').agg({
    'distance': ['sum', 'mean', 'count']
}).reset_index()
yearly_stats.columns = ['year', 'total_distance', 'avg_distance', 'run_count']

# Create bubble chart
bubble_chart = BubbleChart(area=yearly_stats['total_distance'],
    bubble_spacing=0.1)
bubble_chart.collapse()

# Set up the plot
fig, ax = plt.subplots(figsize=(8, 8), subplot_kw=dict(aspect="equal"))

# Generate colors
num_years = len(yearly_stats)
colors = plt.cm.viridis(np.linspace(0, 1, num_years))

# Prepare data for plotting
plot_data = list(zip(yearly_stats['year'],
                    yearly_stats['total_distance'],
                    yearly_stats['avg_distance']))


# Plot the bubble chart
bubble_chart.plot(ax, plot_data, colors)

# Customize the plot
ax.axis("off")
ax.relim()
ax.autoscale_view()
ax.set_title('Yearly Running Summary', fontsize=16, fontweight='bold',
    fontname='Arial')

plt.tight_layout()
plt.show()

```

Yearly Running Summary



5.0.2 Yearly Running Summary Insights

From 2016 to 2018 there was an obvious gradual increase in volume however, 2019/2020 showed a drastic reduction and completely missing from 2021/2022. This is likely due to **COVID-19**

Pandemic and has inadvertently affected fitness overall.

```
[22]: latest_df.weekday.value_counts()

# I really have never ran a monday in 8 weeks lol, I can safely say this is_
↳ because Sundays are long run days so naturally
# Mondays are not really an option
```


```
[22]: weekday
2      11
6       7
4       6
5       5
3       3
1       3
Name: count, dtype: int64
```

```
[23]: # First, let's create a mapping of weekday numbers to names
weekday_names = {
    0: 'Monday',
    1: 'Tuesday',
    2: 'Wednesday',
    3: 'Thursday',
    4: 'Friday',
    5: 'Saturday',
    6: 'Sunday'
}
```

```
[24]: # Calculate the mean moving time for each weekday
result = latest_df.groupby('weekday').agg({
    'distance': 'mean'
}).reindex(range(7)) # This ensures all weekdays are included

# Rename the index with weekday names
result.index = result.index.map(weekday_names)

# Plot the results
ax = result.plot(kind='bar', figsize=(8, 5))
plt.title('Average Distance by Weekday')
plt.xlabel('Weekday')
plt.ylabel('Average Distance')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



5.0.3 Average Distance by Weekday Insights

Weekends (and fridays) are typically when I have the most freedom to run long distances since the availability for training is greater, which also contributes to most of the weekly volume.

```
[25]: fig = plt.figure(figsize=(10, 6))
ax1 = fig.add_subplot(111)

x = np.asarray(latest_df.start_date)
y = np.asarray(latest_df.pace)

ax1.scatter(x, y)
ax1.set_title('Pace over Time')

x2 = mdates.date2num(x)
z = np.polyfit(x2, y, 1)
p = np.poly1d(z)

# Extract the gradient (slope)
gradient = z[0]

plt.plot(x, p(x2), 'r--', linewidth=2)

# Add gradient text to the plot
ax1.text(0.05, 0.95, f'Gradient: {gradient:.4f}',
```

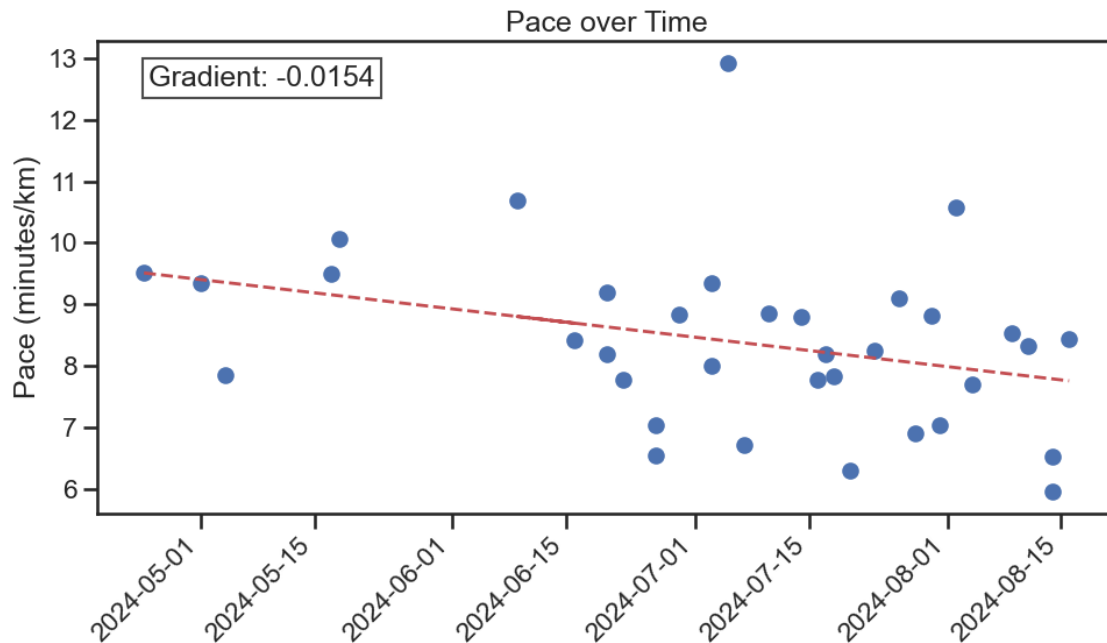
```

transform=ax1.transAxes,
↪verticalalignment='top',bbox=dict(facecolor='white',edgecolor='black',alpha=0.
↪7))

ax1.set_ylabel('Pace (minutes/km)')
# ax1.legend()

fig.autofmt_xdate(rotation=45)
fig.tight_layout()
plt.show()

```



5.0.4 Pace over Time Insights

The graph depicts a slight overall improvement in speed (lower pace) despite day-to-day variations. In my case, its worth mentioning that this is the result of training and all the effort dedicated to the sport.


6 Feature Selection

This is an important step for the application of machine learning models later. This way we can confirm whether the empirically cited characteristics are really good.

The first way to understand the features and how they interact with the response variable is through the `correlation matrix`.

```
[26]: numerical_cols = latest_df.select_dtypes(include=['float64','int32'])

corr = numerical_cols.corr()
plt.figure(figsize = (12,8))
sns.heatmap(corr, fmt=".2f");
plt.title('Correlation between dataset variables')
plt.show()
```



6.0.1 Matrix Analysis

From this matrix, it can be defined that the values closest to +1 have a positive correlation and the values closest to -1 have a negative correlation.

To continue with the analyses, the running dataframe will be scrambled so that problems related to bias and dataframe sequence learning are prevented.

```
[32]: def get_best_rfe_features(X,y, model):
    rfe = RFE(model, step=0.05).fit(X, y)
    selected_features = [i for i, j in zip(X.columns, rfe.support_) if j]
    return selected_features
```

```

runs = numerical_cols.sample(frac=1).reset_index(drop=True)
# list_nulls(runs)

y = runs['pace']
X = runs.drop('pace',axis=1)

encoded_y = preprocessing.LabelEncoder().fit_transform(y)
model = LinearRegression()
linear_feats = get_best_rfe_features(X, encoded_y, model)
linear_feats

```

```

[32]: ['distance',
       'elevation_loss',
       'average_grade',
       'average_heart_rate',
       'month',
       'moving_time_minutes',
       'max_pace']

```

6.0.2 Recursive Feature Elimination

The above code will train the RFE model and recursively eliminate features according to their lack of importance.

Finally `linear_feats` then are the best variables selected for regression

7 Machine Learning

I will only be using 1 machine learning techniques, among the many available, which is **Clustering**. This is because I have no need for regression to predict when I can easily calculate the pace/distance with the right metrics. It is more important to me what kinds of run I'm having and the kinds of heart rate I should be aiming for.

7.0.1 Clustering

I will be performing a technique called K-means. This technique will group records of races that are similar. Note that the method is stochastic, so each execution will likely generate different results.

```

[28]: X = runs.drop('pace',axis=1)
      X = pd.get_dummies(X)

      model = KMeans(n_clusters=4).fit(X)
      clustering_runs = runs.copy()
      clustering_runs['Cluster'] = model.labels_

```

```

[29]: clustering_runs['Cluster'].value_counts()

```

```
[29]: Cluster
      2    16
      0    12
      1     6
      3     1
      Name: count, dtype: int64
```

```
[30]: clustering_runs.groupby('Cluster').mean()
```

```
[30]:      moving_time  distance  elevation_gain  elevation_loss \
Cluster
0      1195.166667    2.474167         63.215174         67.928377
1      4947.000000    9.001667        404.646746        403.980249
2      2449.437500    5.008750        177.150003        175.712498
3      8935.000000   14.810000       992.799927       1008.799988

      elevation_low  elevation_high  max_grade  average_grade \
Cluster
0          42.825000         67.708333  31.144058         -0.246359
1         -16.466667         76.350000  46.251586          0.028638
2          53.287499        103.237501  37.202141          0.035288
3          395.200012        523.000000  49.956367         -0.156001

      max_heart_rate  average_heart_rate  total_steps  year  month \
Cluster
0          177.333333         162.641992    3004.833333  2024.0  6.750000
1          175.273810         156.252586   11943.000000  2024.0  6.833333
2          182.437500         165.680799    6201.375000  2024.0  6.625000
3          202.000000         181.586426   20382.000000  2024.0  5.000000

      moving_time_minutes      pace  max_pace
Cluster
0          19.919167    7.996667  4.294167
1          82.450000    9.328333  3.453333
2          40.824375    8.238125  3.228750
3         148.920000   10.060000  4.980000
```

7.1 Cluster Overview

There are 4 distinct clusters, each representing a different type of run or running pattern

7.1.1 Clustering Characteristics

*Cluster A: **Moderate** Runs* - Average distance: ~5km - Moderate elevation gain - Higher heart rates - Representing typical training runs (also the highest sample among the clusters)

*Cluster B: **Long** Runs* - Average distance: ~9km - Higher elevation gain - Lower average heart rate than Moderate Run (likely due to long runs being used for slower, volume jogs) - Endurance building runs

*Cluster C: **Short** Runs* - Average distance: ~2.5km - Lowest elevation gain, flattest ground - Higher pace than the rest - Heart rate roughly matching Cluster 0 even with less distance - Representing speed sessions and/or training sessions

*Cluster D: **Challenging** Long Runs* - Average distance: ~15km - Highest elevation gain - Highest heart rate - Outlier, representing races and trail runs

7.1.2 Interesting Observations:

Run distance The longer runs(Cluster 1) have lower average heart rate than shorter runs (Cluster 2), suggesting better cardiovascular efficiency on longer distances. #### Elevation vs Pace Clusters with higher elevation gain/loss (1 and 3) have slower paces, showing the impact of terrain on speed.

7.1.3 Potential Insights:

Training Variety: A good mix of run types, from short, intense sessions to longer endurance runs. *Heart Rate Zones:* Heart rate varies predictably with run intensity and distance. *Terrain Impact:* There's a clear relationship between elevation changes and pace/effort.

7.2 Calculate the race!

Now to conclude this study, lets predict how long it would take for me to run a 21KM run this November.

Lets try and compare 2 ways of calculating,

1. one by (naively) looking at the gradient of improvement over the past 8 weeks vs
2. one by looking at the rolling average as well as extra weightage on longer runs

7.2.1 –Gradient–

```
[46]: # Calculate the improvement rate
days = (latest_df.start_date.max() - latest_df.start_date.min()).days
improvement_rate = gradient * days # Improvement in minutes/km over the entire
    ↪ period

# Project future performance
target_pace = 6.0 # 2:06:00 for a half marathon is about 6 min/km
# Get the most recent date
most_recent_date = latest_df['week_start_date'].max()

# Calculate the date two weeks ago
two_weeks_ago = most_recent_date - timedelta(weeks=2)

# Filter the DataFrame for the last two weeks
last_two_weeks_df = latest_df[latest_df['week_start_date'] > two_weeks_ago]

# Get the highest pace from the last two weeks
current_pace = last_two_weeks_df['pace'].mean()
```

```

time_to_target = (current_pace - target_pace) / (-gradient)

def days_to_months_days(total_days):
    average_days_per_month = 30

    # Calculate the number of months
    months = int(total_days // average_days_per_month)

    # Calculate the remaining days
    remaining_days = int(total_days % average_days_per_month)

    return months, remaining_days

months, days = days_to_months_days(time_to_target)

print(f"Current pace: {current_pace:.2f} min/km")
print(f"Improvement rate: {-gradient*365:.2f} min/km per year")
print(f"Estimated days to reach target pace: {time_to_target:.0f}")
print(f"Estimated months and days to reach target pace: {months} months, {days:.0f} days")

```

Current pace: 7.55 min/km

Improvement rate: 5.62 min/km per year

Estimated days to reach target pace: 101

Estimated months and days to reach target pace: 3 months, 10 days

7.2.2 –Rolling Average–

```

[45]: def predict_time_to_goal_pace(df, goal_pace, weeks_for_average=4,
    ↪long_run_threshold=10, long_run_weight=1.5):
    # Sort the dataframe by date
    df = df.sort_values('activity_date')

    # Weight long runs more heavily
    weighted_pace = pd.Series(np.where(df['distance'] >= long_run_threshold,
    df['pace'] * long_run_weight,
    df['pace']))

    # Calculate weighted rolling average
    weighted_rolling_avg = weighted_pace.rolling(window=7*weeks_for_average,
    ↪min_periods=1).mean()

    # Calculate the rate of improvement (pace decrease per day)
    days = (df['activity_date'] - df['activity_date'].min()).dt.days
    improvement_rate = np.polyfit(days, weighted_rolling_avg, 1)[0]

```

```

# Get the most recent weighted rolling average
current_weighted_avg = weighted_rolling_avg.iloc[-1]

# Estimate current half marathon pace (typically 5-10% slower than average
↳training pace)
current_hm_pace = current_weighted_avg * 1.07 # Assuming 7% here

# Calculate time to reach goal pace
time_to_goal = (current_hm_pace - goal_pace) / (-improvement_rate)

return current_hm_pace, improvement_rate, time_to_goal

# Example usage
goal_pace = 6.0 # min/km
current_pace, improvement_rate, days_to_goal =
↳predict_time_to_goal_pace(latest_df, goal_pace)

print(f"Current estimated Half Marathon Pace: {current_pace:.2f} min/km")
print(f"Current improvement rate: {-improvement_rate*365:.4f} min/km per year")
print(f"Estimated days to reach goal pace: {days_to_goal:.0f}")

# Convert days to months and remaining days
months = int(days_to_goal // 30)
remaining_days = int(days_to_goal % 30)
print(f"Estimated time to reach goal pace: {months} months and {remaining_days}
↳days")

```

```

Current estimated Half Marathon Pace: 9.24 min/km
Current improvement rate: 2.5660 min/km per year
Estimated days to reach goal pace: 460
Estimated time to reach goal pace: 15 months and 10 days

```

8 Conclusion

The comparison of two prediction methods for the upcoming 21KM run in November reveals significantly different outcomes. The **gradient** method, a simpler approach, suggests a current pace of 7.55 min/km and predicts reaching your goal in about 3 months. This method paints an **optimistic picture of rapid improvement**.

In contrast, the **rolling average method**, which factors in longer runs, estimates a slower current pace of 9.24 min/km and projects a longer timeframe of 15 months to reach your goal. This approach likely provides a more realistic prediction, **accounting for the challenges of maintaining pace over longer distances**.

For the November race, it's prudent to base your expectations on the more conservative rolling average estimate while working to exceed it. This balances optimism with realism, setting you up for a potentially satisfying performance without risking disappointment from overly ambitious goals.

Thanks for reading!